



Beginner's Crash Course to Elastic Stack Series

Part 1.2: Understanding the Relevance of your search using
Elasticsearch and Kibana

Lisa Jung
Developer Advocate @Elastic



Beginner's crash course to the Elastic Stack Series

- **Part 1.1: Intro to Elasticsearch and Kibana**
 - use case of Elasticsearch and Kibana
 - the basic architecture of Elasticsearch
 - perform CRUD(Create, Read, Update, and Delete) operations with Elasticsearch and Kibana

Missed the first workshop? No worries!

- **Part 1.1: Intro to Elasticsearch and Kibana**
 - Repo: <https://ela.st/workshop-1-repo>

The Elastic Stack

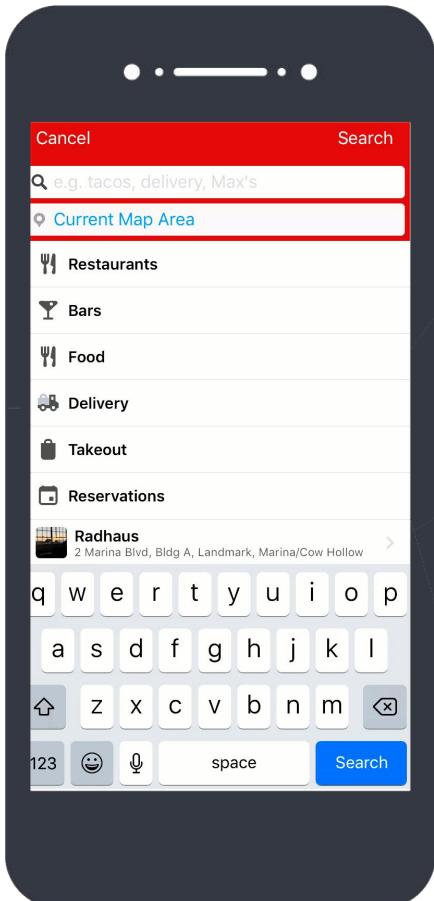
Reliably and securely take data from any source, in any format, then search, analyze, and visualize it in real time.



Elasticsearch

Store | Search | Analyze



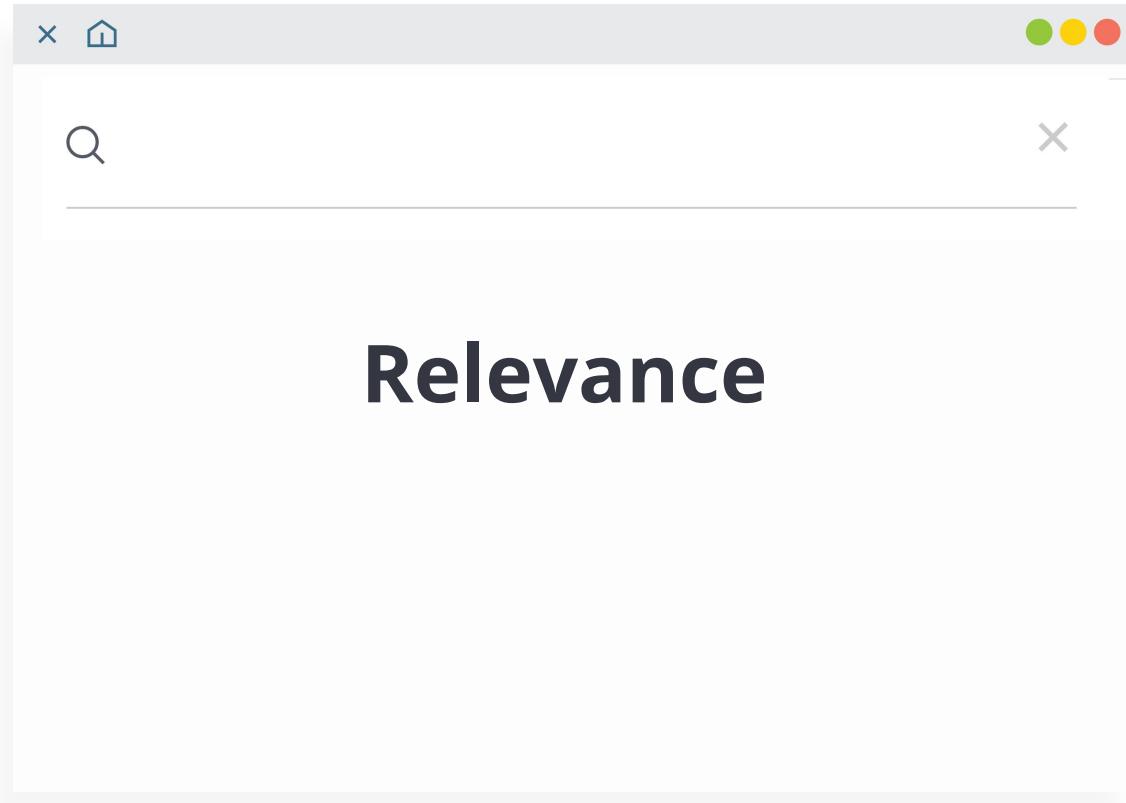


The image shows the Instacart mobile application interface. At the top, there is a navigation bar with a "X" button, a house icon, and three colored dots (green, yellow, red) on the right. The main header features the Instacart logo and a "Stores" dropdown. On the right side of the header are links for "Delivery in 94086", "Account", "Help", and a green "Cart" button with a red notification badge showing the number "4". The background of the screen is a photograph of fresh produce including avocados and kale. In the center, there is a Safeway logo with the word "Safeway" below it and links to "View pricing policy" and "More info". A search bar at the bottom has the prefix "can|". A dropdown menu is open, listing categories under the heading "Canned Goods": "Department", "Canned Fruit & Applesauce", "Canned & Jarred Vegetables", and "Canned Meals & Beans". Below this, there is a "Coupon saving" section with a "Shop Coupons" button, and a "Save Now" button next to a Kraft logo. At the bottom of the screen, there is a message "Based on your cart" and a "View more" link.

Speed, Scale, **Relevance**

Elastic is a search company.

We focus on value to users by producing fast results that operate at scale and are relevant. This is our DNA. We believe search is an experience. It is what defines us, and makes us unique.



How do we measure relevance?

- **Precision**
- **Recall**

Elasticsearch

Store | Search | Analyze



I store data as documents!

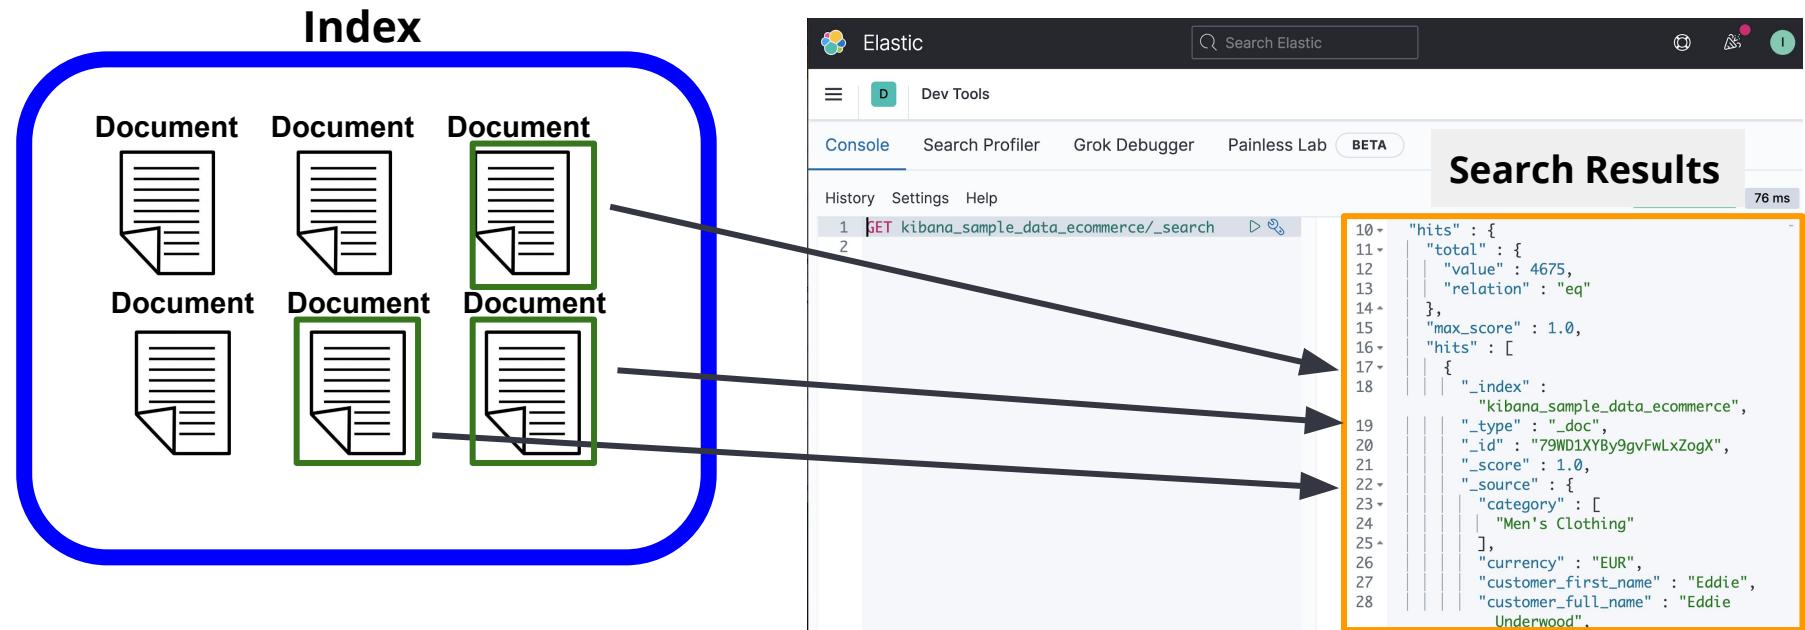
Index

Document Document Document
Document Document Document



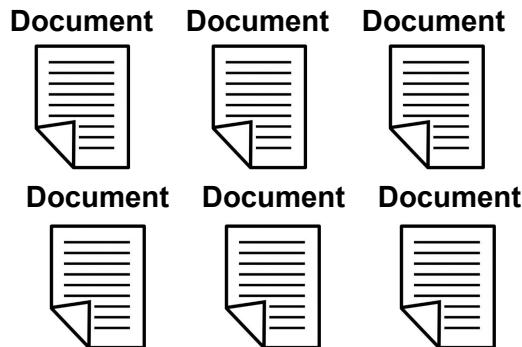
Documents with similar traits are grouped into an index!

When search query is sent, Elasticsearch retrieves relevant documents and presents the documents as search results.

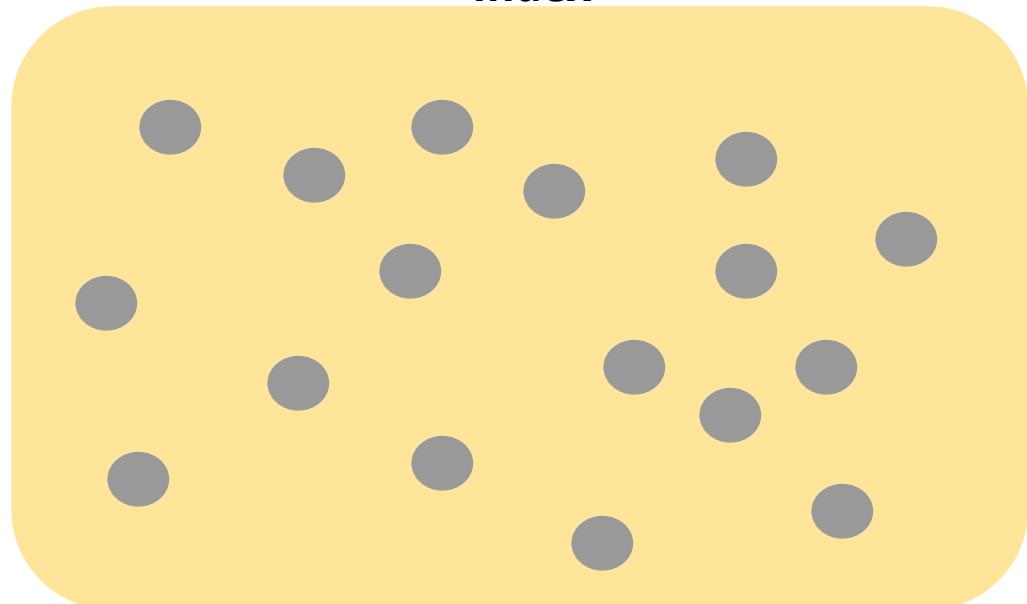


These two diagrams depict the same thing!

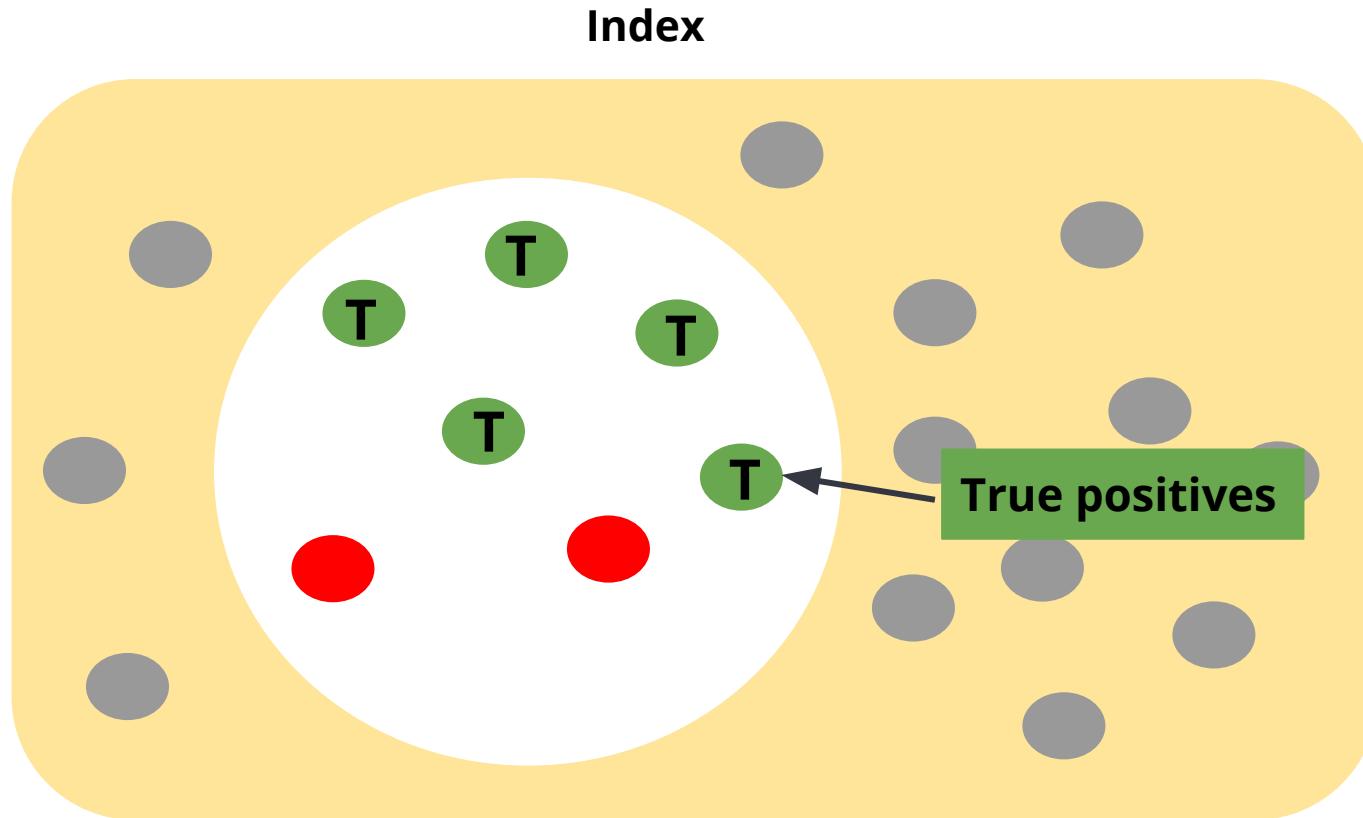
Index



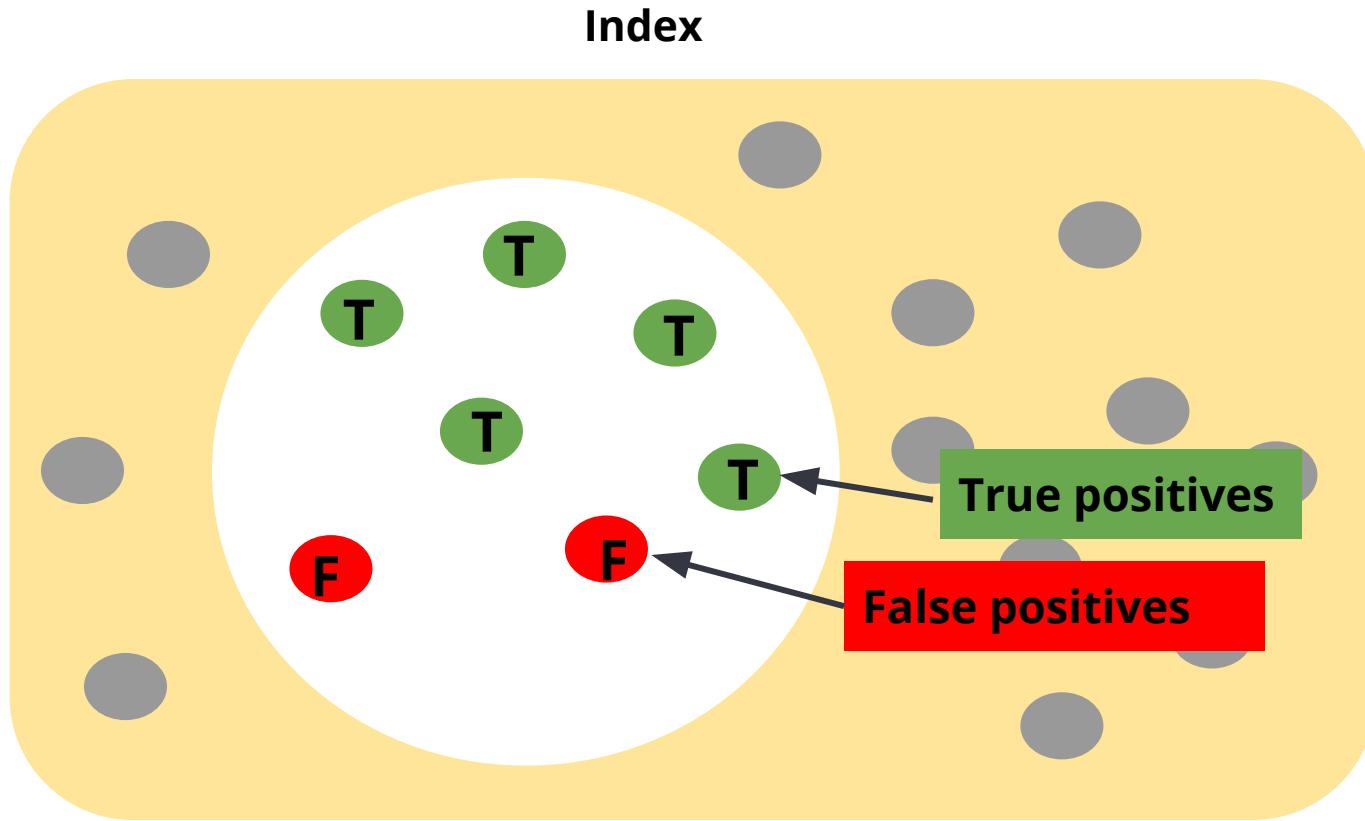
Index



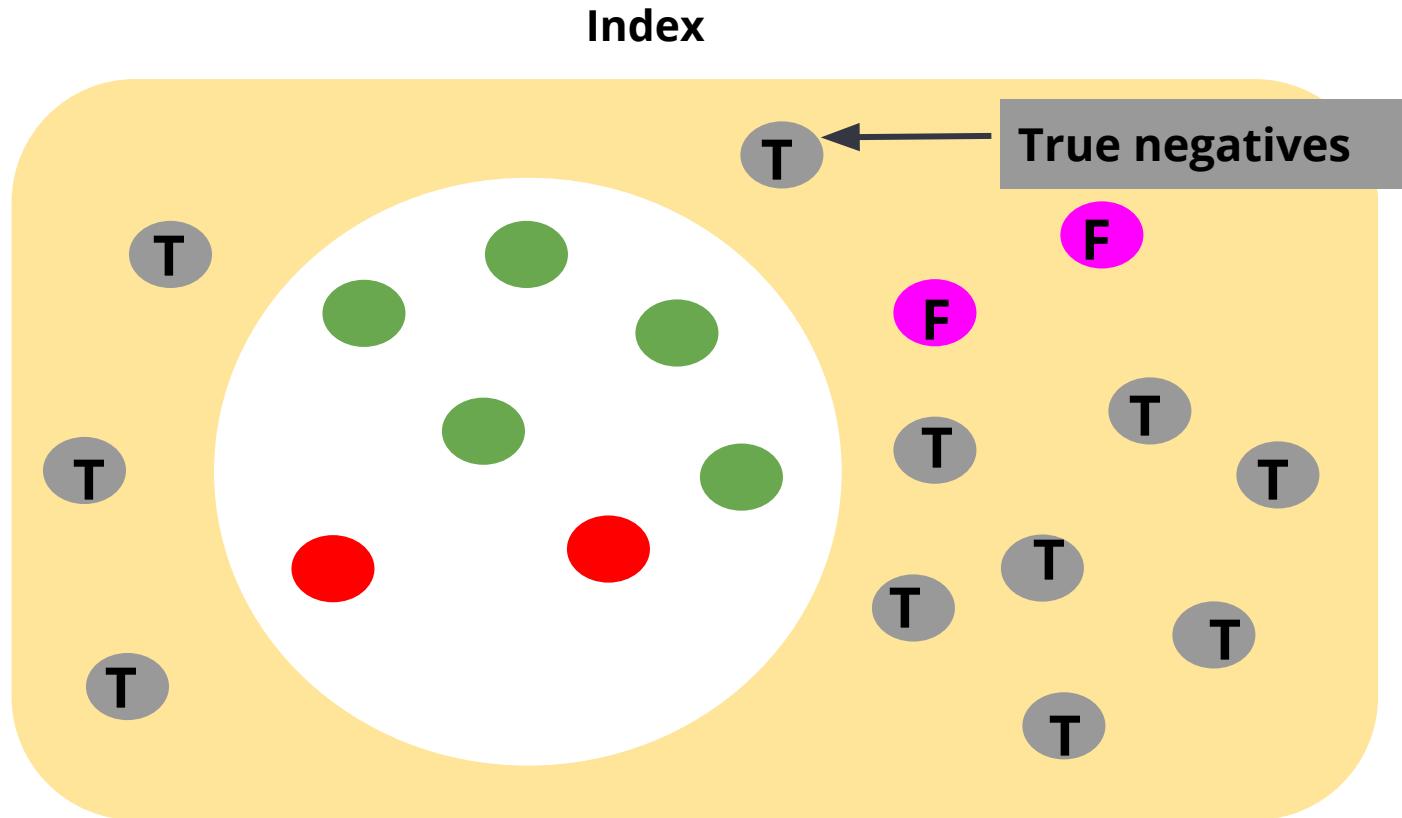
True positives are relevant documents that are returned to the user.



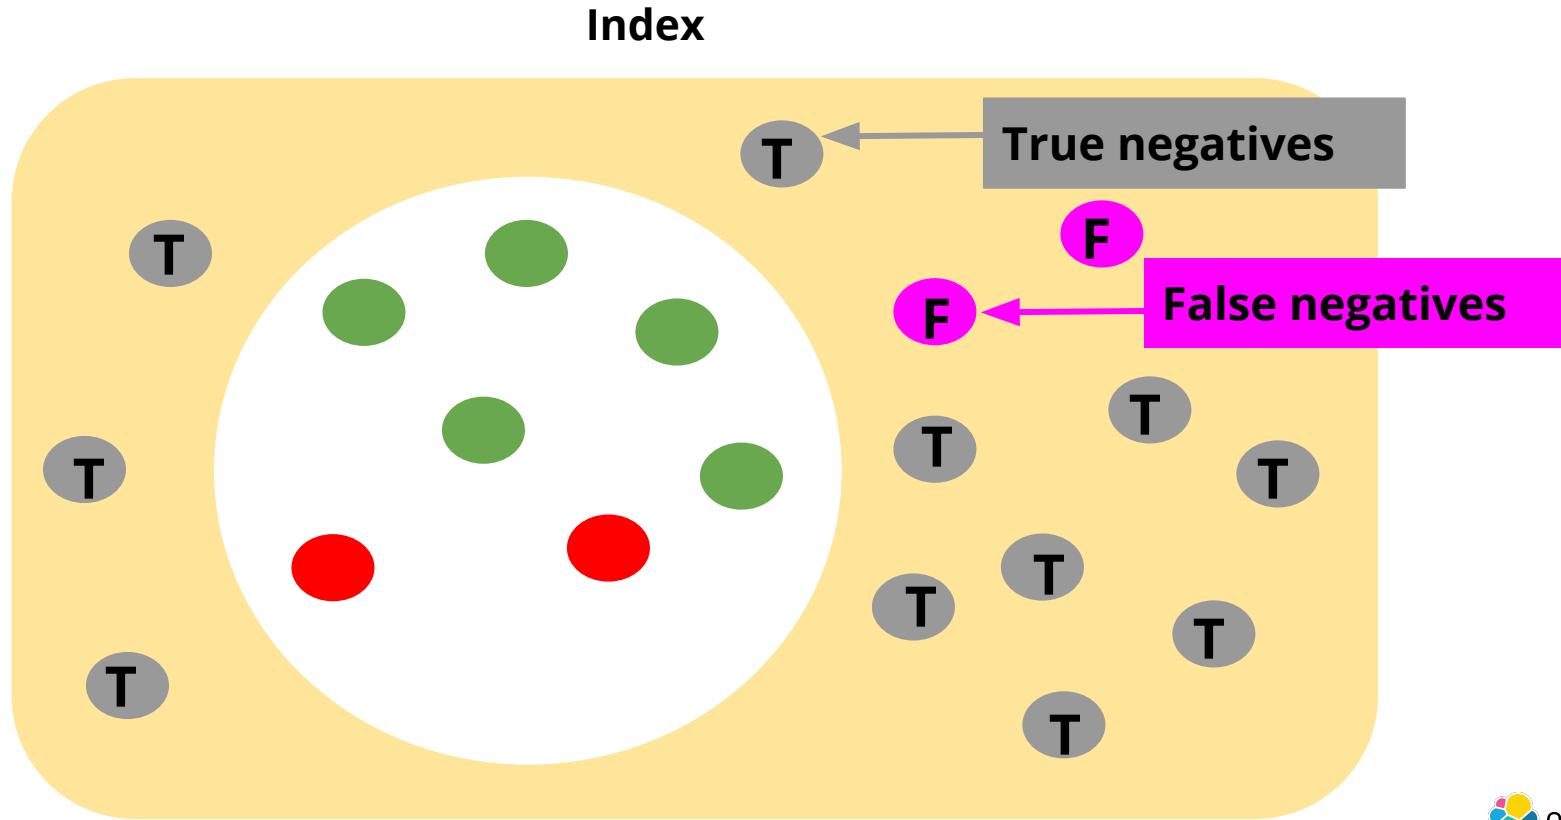
False positives are irrelevant documents that are returned to the user.



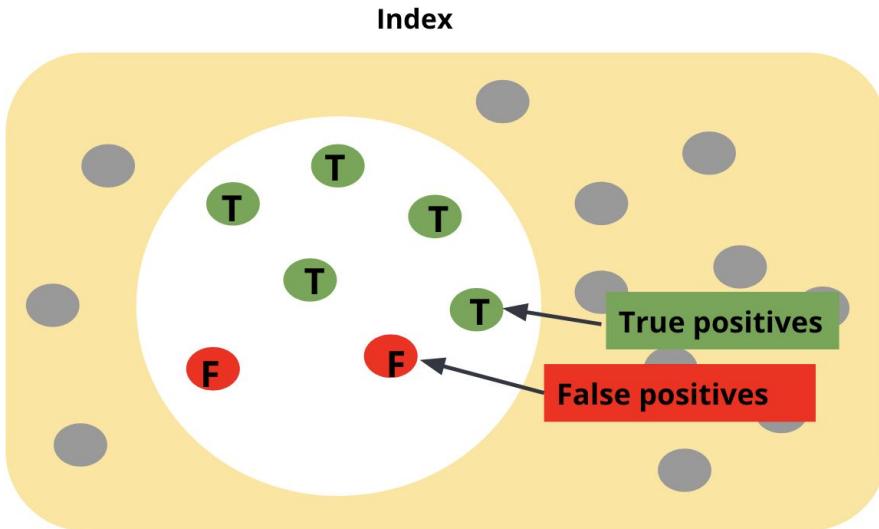
True negatives are irrelevant documents that are not returned to the user.



False negatives are relevant documents that were not returned to the user.



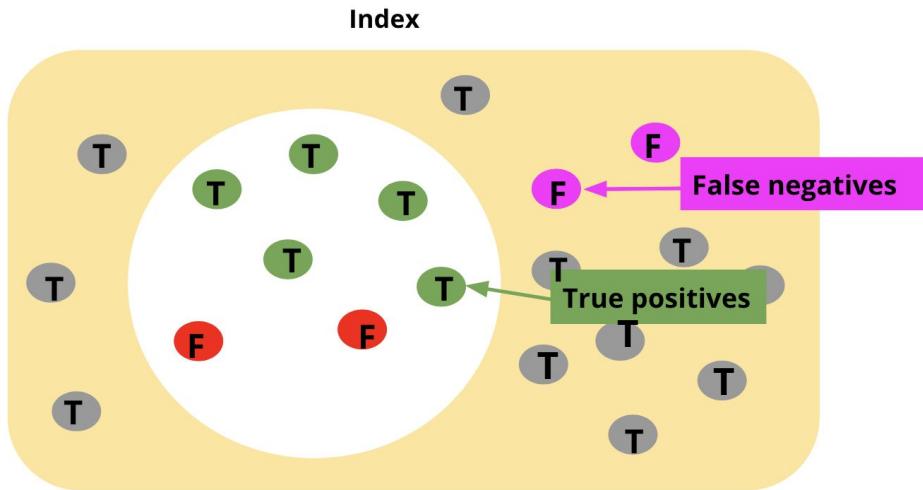
What is precision?



$$\text{Precision} = \frac{\text{True positives}}{\text{True positives} + \text{False positives}}$$

What portion of the retrieved data is actually relevant to the search query?

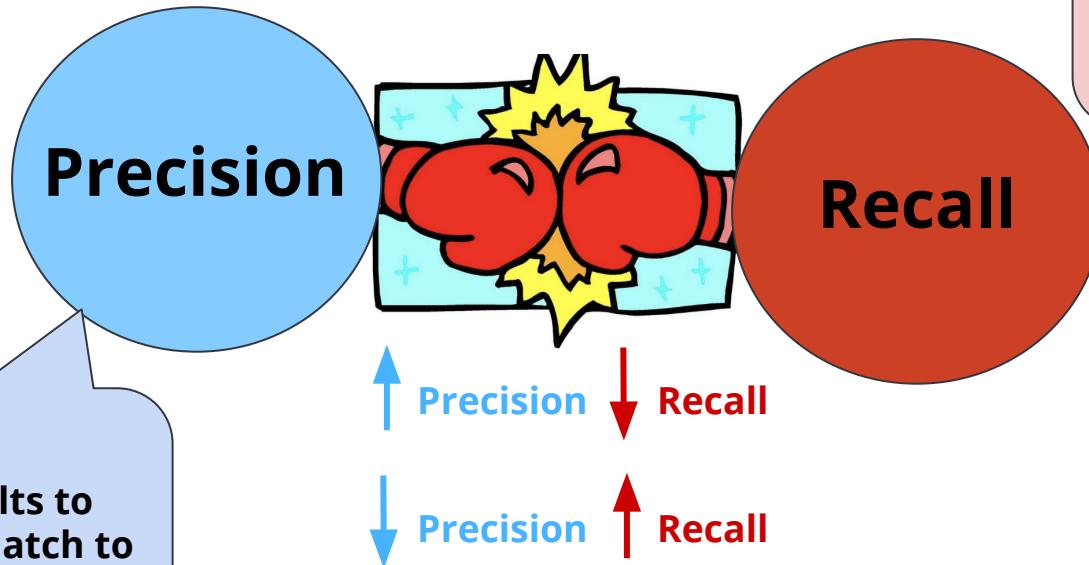
What is recall?



$$\text{Recall} = \frac{\text{True positives}}{\text{True positives} + \text{False negatives}}$$

What portion of relevant data is being returned as search results?

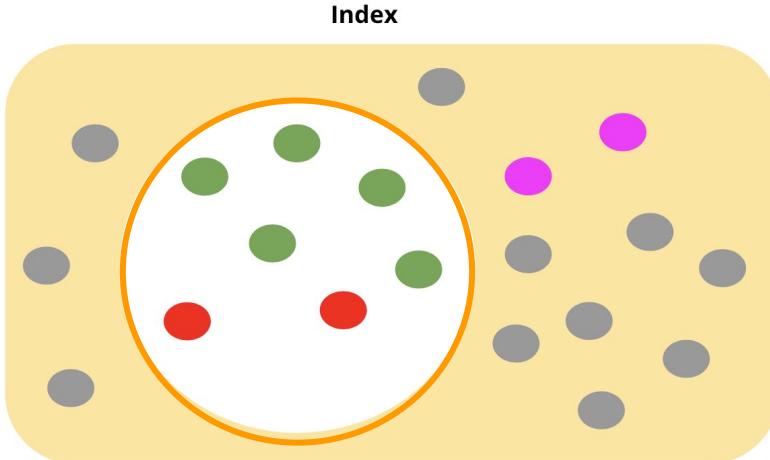
Precision and Recall are inversely related



I want all the retrieved results to be a perfect match to the query, even if it means returning less documents.

I want to retrieve more results even if documents may not be a perfect match to the query.

Precision and recall determine which documents are included in the search results.



Elastic Dev Tools

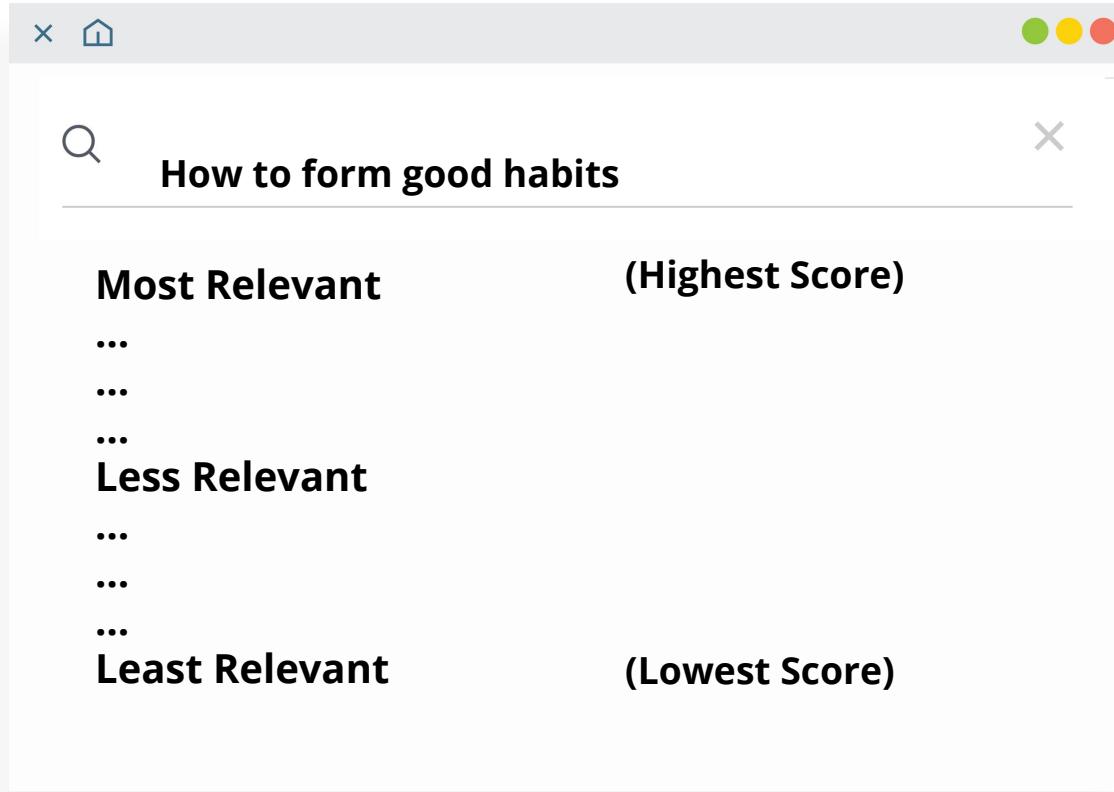
Console Search Profiler Grok Debugger Painless Lab BETA

History Settings Help

```
1 GET kibana_sample_data_ecommerce/_search
2
3
4
5
6
7
8
9
10 "hits" : {
11   "total" : {
12     "value" : 4675,
13     "relation" : "eq"
14   },
15   "max_score" : 1.0,
16   "hits" : [
17     {
18       "_index" :
19         "kibana_sample_data_ecommerce",
20       "_type" : ".doc",
21       "_id" : "79WD1XYBy9gvFwLxZogX",
22       "_score" : 1.0,
23       "_source" : {
24         "category" : [
25           "Men's Clothing"
26         ],
27         "currency" : "EUR",
28         "customer_first_name" : "Eddie",
29         "customer_full_name" : "Eddie Underwood"
30       }
31     }
32   ]
33 }
```

Precision and recall do not determine which of the returned documents are more relevant compared to the other!

Ranking refers to ordering of the results (from most relevant results at the top, to least relevant at the bottom).



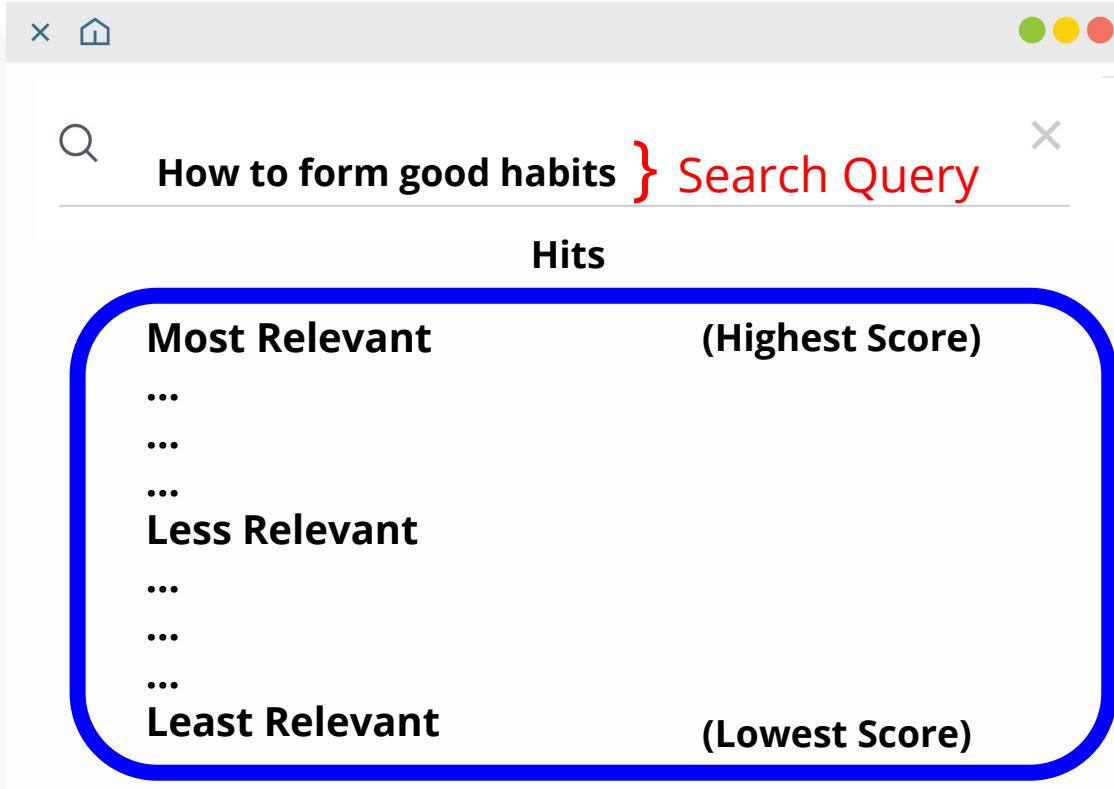
What is score?

- The score is a value that represents how relevant a document is to that specific query
- A score is computed for each document that is a hit

What is score?

- Term Frequency(TF)
- Inverse Document Frequency(IDF)

What is score?



Term Frequency(TF) determines how many times each search term appears in a document.

The screenshot shows a search interface with the following components:

- Search Terms:** "How to form good habits" (with "habits" highlighted in a blue box).
- Search Query:** } Search Query
- Results:** Two JSON documents are listed:
 - Document 1:** "Atomic Habits" by James Clear, self-help category. Description: "No matter your goals, Atomic Habits offers a proven framework for improving every day. James clear, ... habits...habits ...habits" **TF= 4**
 - Document 2:** "The Mental Toughness Handbook" by Damon Zahariades, self-help category. Description: "Imagine boldly facing any challenge that comes your way... 5 daily habits you must embrace to strengthen your mind and harden your resolve. Why willpower and motivation are unreliable..." **TF= 1**

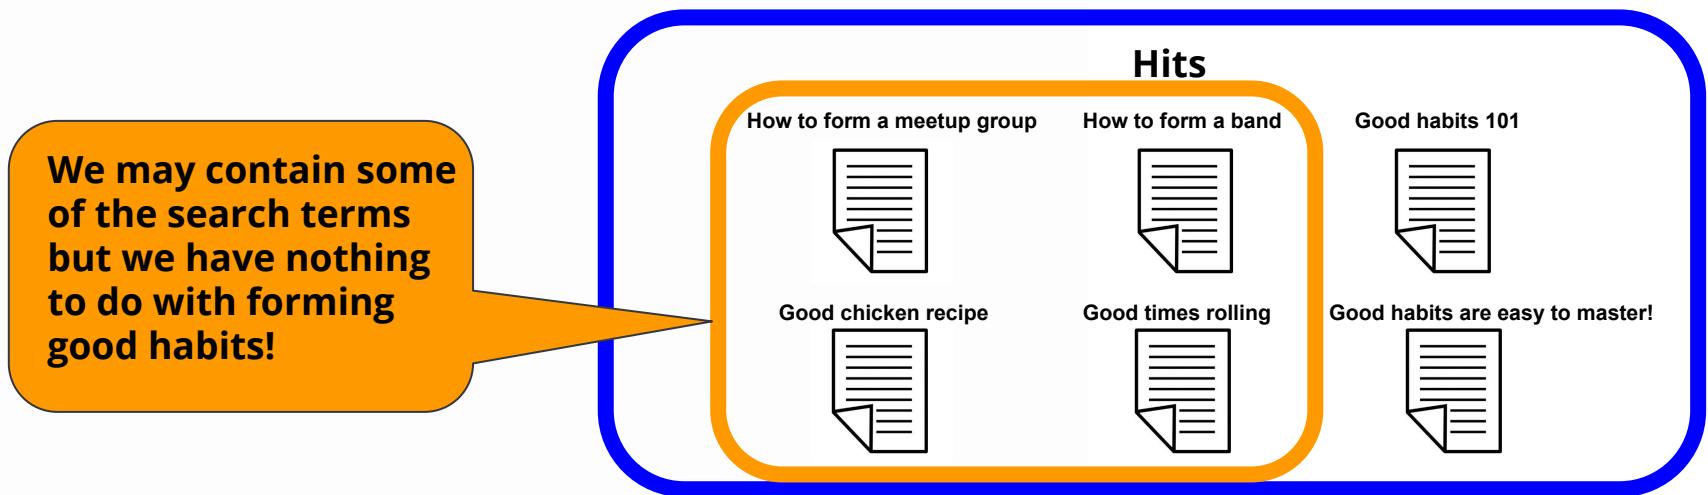
If search terms are found in high frequency in a document, the document is considered more relevant to the search query.

```
{  
  "title": "Atomic Habits",  
  "author": "James Clear",  
  "category": "self-help",  
  "description": "No matter your goals, Atomic Habits offers a proven framework for improving every day. James clear, ... habits...habits ...habits" TF= 4  
}  
  
{  
  "title": "The Mental Toughness Handbook",  
  "author": "Damon Zahariades",  
  "category": "self-help",  
  "description": "Imagine boldly facing any challenge that comes your way... 5 daily habits you must embrace to strengthen your mind and harden your resolve. Why willpower and motivation are unreliable..." TF= 1  
}
```

What is Inverse Document Frequency(IDF)?



IDF diminishes the weight of terms that occur very frequently in the document set and increases the weight of terms that occur rarely!



Fine tuning precision or recall using Elasticsearch and Kibana



Click on the link to the workshop repo.

<https://ela.st/workshop-2-repo>

Scroll down to the Resources section & open Free Elastic Cloud Trial link in a new tab.

Beginner's Crash Course to the Elastic Stack Series

Part 1.2: Understanding the relevance of your search with Elasticsearch and Kibana

Welcome to the Beginner's Crash Course to the Elastic Stack!

This repo contains all resources shared during the workshop 1.2: Understanding the relevance of your search with Elasticsearch and Kibana.

Resources

- [Free Elastic Cloud Trial](#)
 - Open Link in New Tab
 - Open Link in New Window
 - Open Link in Incognito Window
 - Save Link As...
 - Copy Link Address
 - Copy
 - Search Google for "Free Elastic Cloud Trial"
 - Print...
 - LastPass
 - Inspect
 - Speech
 - Services
- [Instructions for download](#)
- [Presentation](#)
- [Data set from Kaggle](#)

Elastic Austin User Group Want to attend live workshops? Join the Elastic Austin User Group to keep up to date on all



Scroll down to the Resources section & open Free Elastic Cloud Trial link in a new tab.

The screenshot shows two main sections. On the left, a blue header reads "ELASTICSEARCH SERVICE" with a logo. Below it, a section titled "Deploy Elasticsearch and Kibana in 3 minutes or less" contains text about a 30-day free trial and a "Start Free Trial" button. A red box highlights the "Enter your email" input field and the "Start Free Trial" button. On the right, a "First cluster" configuration interface shows "Edit Data" for "grp.data.HighIO.1" (A Kibana Instance). It includes tabs for Deployments, Logs, Snapshots, API Console, Kibana, APM, Metrics, Security, Performance, and Custom plugins. The "Fault tolerance" section shows 1 zone selected. RAM per Node is set to 8 GB. The "Architecture" section shows two zones: Zone 1 with 1 instance and Zone 2 with 1 instance, each having 1 GB of RAM per zone. A "Summary" table provides cluster statistics.

Trial comes with the following



8 GB memory



240 GB storage



High availability across two zones



One-click upgrade to the latest versions of Kibana and



Advanced security features like authentication and role-based access control



Alerting capabilities to trigger notifications



Monitoring to optimize your cluster health



Protect your cluster with Encryption at Rest

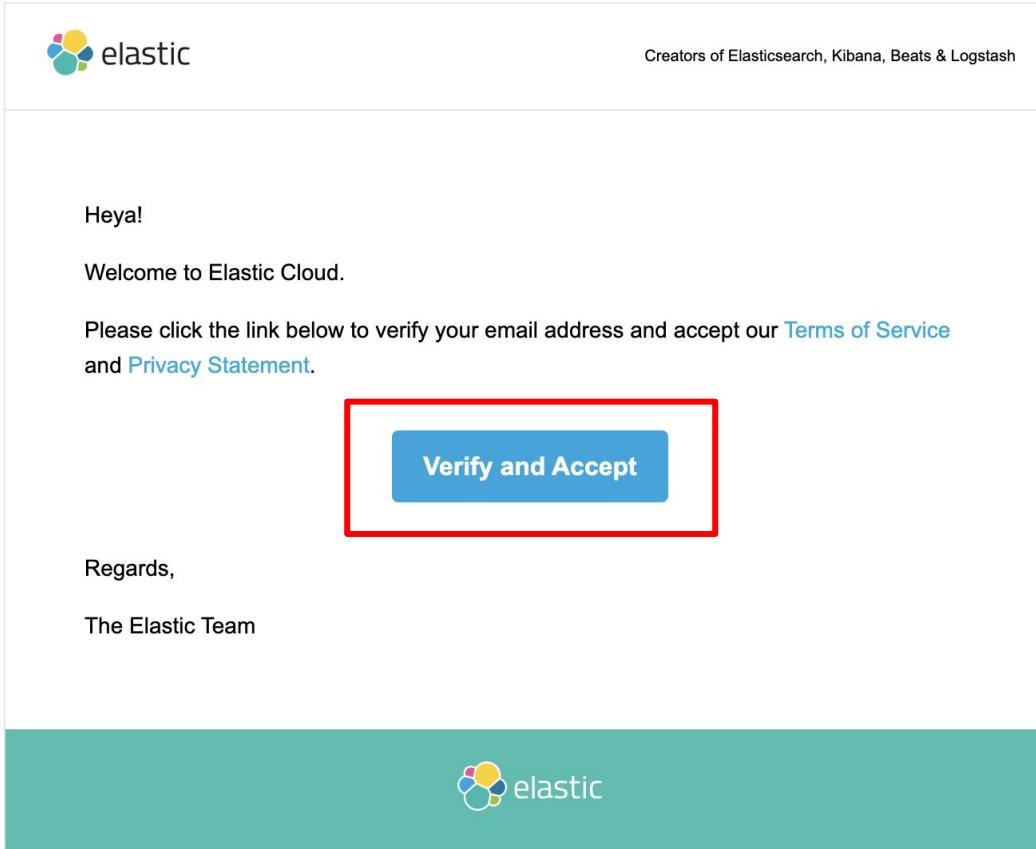
Go to the email account you signed up with and verify your email.



Please Verify Your Email

Splendid! Head to your inbox! You should see an email verifying your account and quick-start information.

Open email from Elastic. Click on verify and accept button.



The image shows a screenshot of an email from the company 'elastic'. The email header includes the 'elastic' logo and the text 'Creators of Elasticsearch, Kibana, Beats & Logstash'. The body of the email starts with 'Heya!', followed by 'Welcome to Elastic Cloud.' and a message asking the recipient to click a link to verify their email address and accept the 'Terms of Service' and 'Privacy Statement'. A blue button labeled 'Verify and Accept' is centered below the message, and it is highlighted with a red rectangular box. The email concludes with 'Regards,' and 'The Elastic Team'. The footer of the email features the 'elastic' logo again.

Heya!

Welcome to Elastic Cloud.

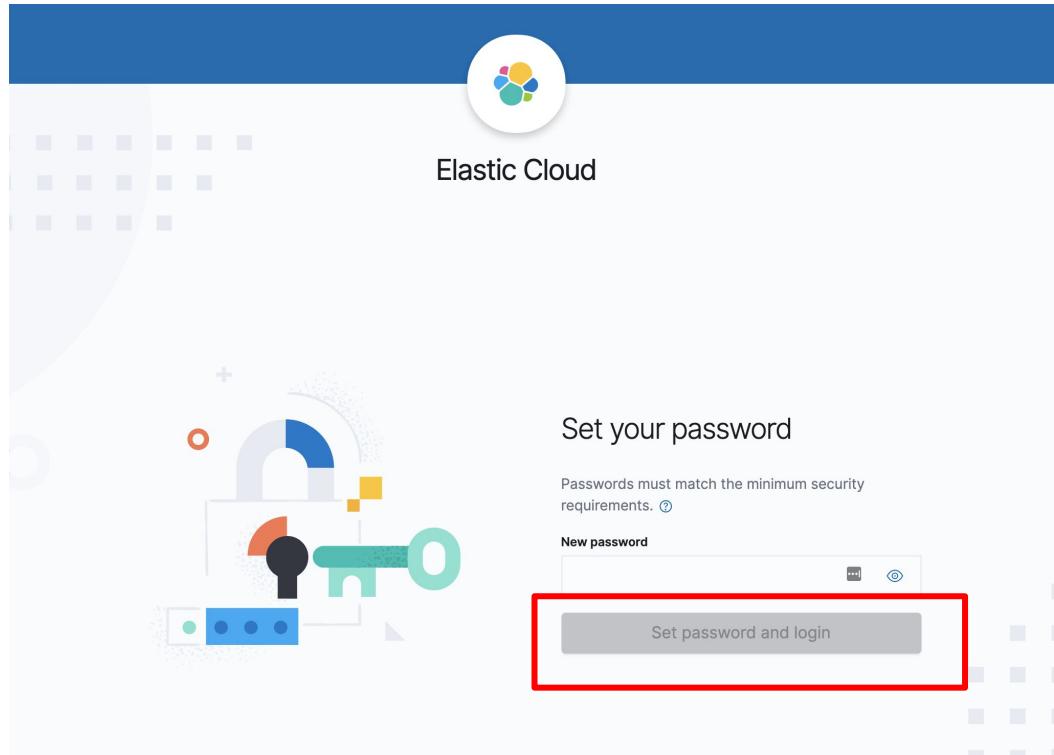
Please click the link below to verify your email address and accept our [Terms of Service](#) and [Privacy Statement](#).

Verify and Accept

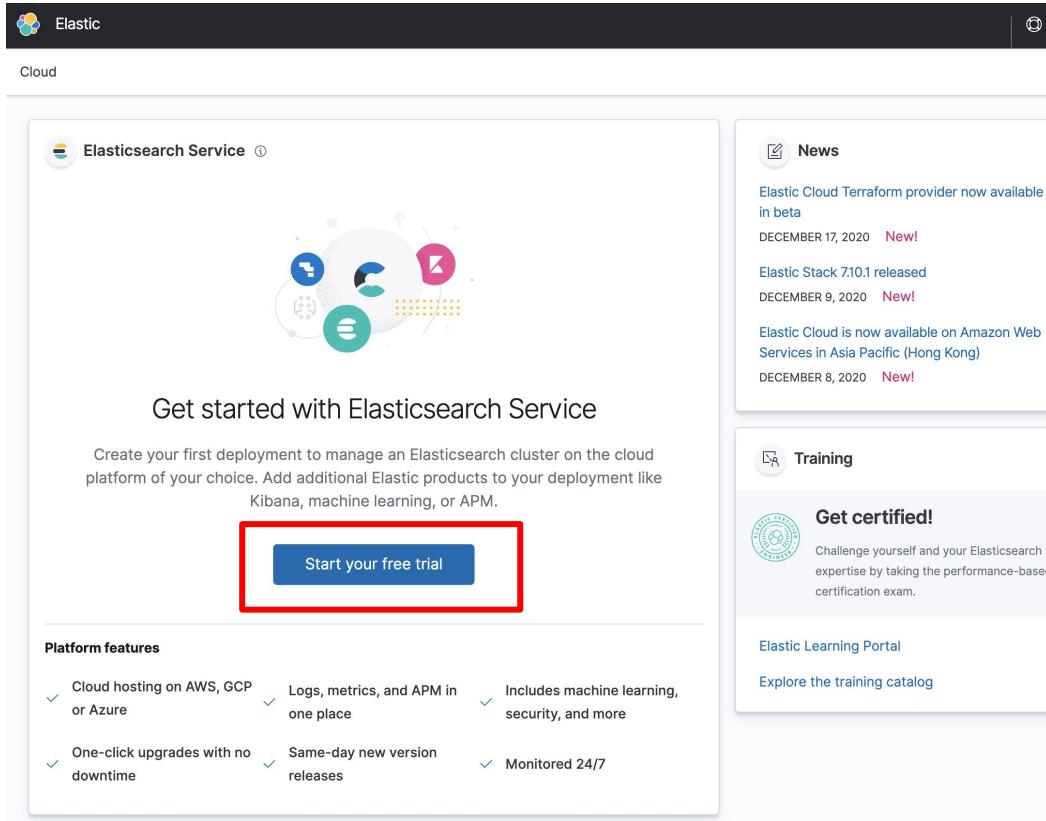
Regards,

The Elastic Team

Enter your password.



Click on start your free trial.



The screenshot shows the Elastic Cloud interface for the Elasticsearch Service. At the top left is the Elastic logo. Below it, a navigation bar includes 'Cloud' and other service icons. The main content area features a large 'Elasticsearch Service' section with a circular icon containing a gear and a magnifying glass. Below this is a heading 'Get started with Elasticsearch Service'. A paragraph describes creating a deployment on the cloud platform of choice, mentioning Kibana, machine learning, or APM. A prominent blue button labeled 'Start your free trial' is centered, with a red rectangular box drawn around it. To the right of this main section are two sidebar boxes: 'News' (listing recent releases like 'Elastic Cloud Terraform provider now available in beta') and 'Training' (with a 'Get certified!' section and a 'Elastic Learning Portal' link).

Elasticsearch Service

Get started with Elasticsearch Service

Create your first deployment to manage an Elasticsearch cluster on the cloud platform of your choice. Add additional Elastic products to your deployment like Kibana, machine learning, or APM.

[Start your free trial](#)

Platform features

- ✓ Cloud hosting on AWS, GCP or Azure
- ✓ Logs, metrics, and APM in one place
- ✓ Includes machine learning, security, and more
- ✓ One-click upgrades with no downtime
- ✓ Same-day new version releases
- ✓ Monitored 24/7

News

Elastic Cloud Terraform provider now available in beta

DECEMBER 17, 2020 **New!**

Elastic Stack 7.10.1 released

DECEMBER 9, 2020 **New!**

Elastic Cloud is now available on Amazon Web Services in Asia Pacific (Hong Kong)

DECEMBER 8, 2020 **New!**

Training

Get certified!

Challenge yourself and your Elasticsearch expertise by taking the performance-based certification exam.

[Elastic Learning Portal](#)

[Explore the training catalog](#)

Select the Elastic Stack.

The screenshot shows the Elasticsearch Service interface for creating a deployment. The left sidebar includes links for Deployments, Create deployment, Extensions, API keys, Traffic filters, and Help. The main content area features a welcome message for a 30-day free trial and a section asking what you want to do with your data. Four service options are listed:

- Elastic Stack**: Deploy Elasticsearch and Kibana. Search, analyze, and visualize data from any source, in any format. **Select** button highlighted with a red box.
- Elastic Enterprise Search**: Easily implement powerful search experiences for your website, app, or workplace with refined APIs and tools. **Select** button.
- Elastic Observability**: Consolidate your logs, metrics, application traces, and system availability with purpose-built UIs. **Select** button.
- Elastic Security**: Prevent, collect, detect, and respond to threats for unified protection across your infrastructure. **Select** button.

Configure your settings.

Select hardware profile

I/O Optimized Recommended

Use for all-purpose workloads, including time-series data like logs and metrics. [See details](#)

Compute Optimized

Run CPU-intensive workloads or run smaller workloads cost-effectively when you need less memory and storage. [See details](#)

Memory Optimized

Perform memory-intensive operations efficiently, including workloads with frequent aggregations. [See details](#)

Hot-Warm Architecture

Useful for time-series analytics that benefit from automatic index curation. [See details](#)

Cross Cluster Search Not available in trial

Search data across one or more associated remote deployments. [See details](#)

Deployment settings

Choose the cloud provider, region, and Elastic Stack version.

Cloud provider

Pick a cloud and let us handle the rest. No additional accounts required.

 Google Cloud  Azure  Amazon Web Services

Region

Select the location of your deployment.

 West US 2 (Washington)

Version

Choose the Elastic Stack version.

7.10.1

Name your deployment then create deployment.

Region

Select the location of your deployment.

 West US 2 (Washington)

Version

Choose the Elastic Stack version.

7.10.1

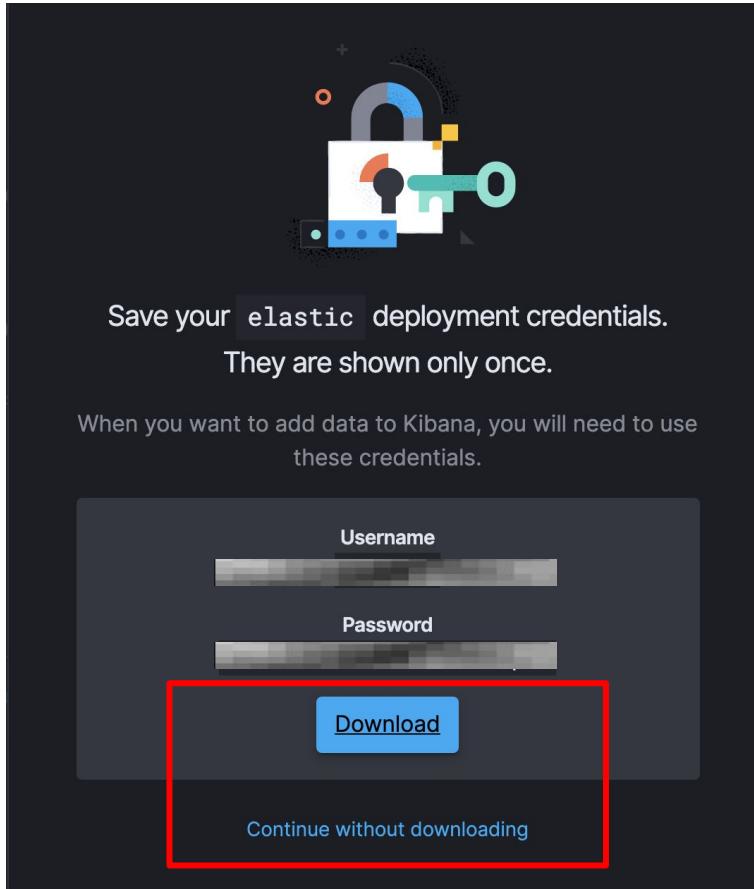
Name your deployment

You can always change this later.

Beginner's Crash Course

 Create deployment

Save your deployment credentials.



Open Kibana.

← → ⌛ https://cloud.elastic.co/deployments/316be9b6fe7c451cba113e1b2cfe35c5

Elastic

Cloud | Deployments | Beginner's Crash Course to the Elastic Stack

Deployments

Beginner's Crash Course to t...
Edit
Elasticsearch
Snapshots
API console
Kibana
APM
Enterprise Search
Logs and metrics
Activity
Security
Performance

Extensions

API keys

Traffic filters

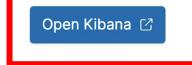
Help

Beginner's Crash Course

us-central1 (Iowa)

Get started with your deployment

The next step is to ingest data and create visualizations in Kibana.

Open Kibana 

Forgot to save your credentials?
[Reset your deployment password](#)

Deployment name: Beginner's Crash Course [Edit](#)
Deployment ID: 316be9b6fe7c451cba113e1b2cfe35c5

Deployment status:  Healthy

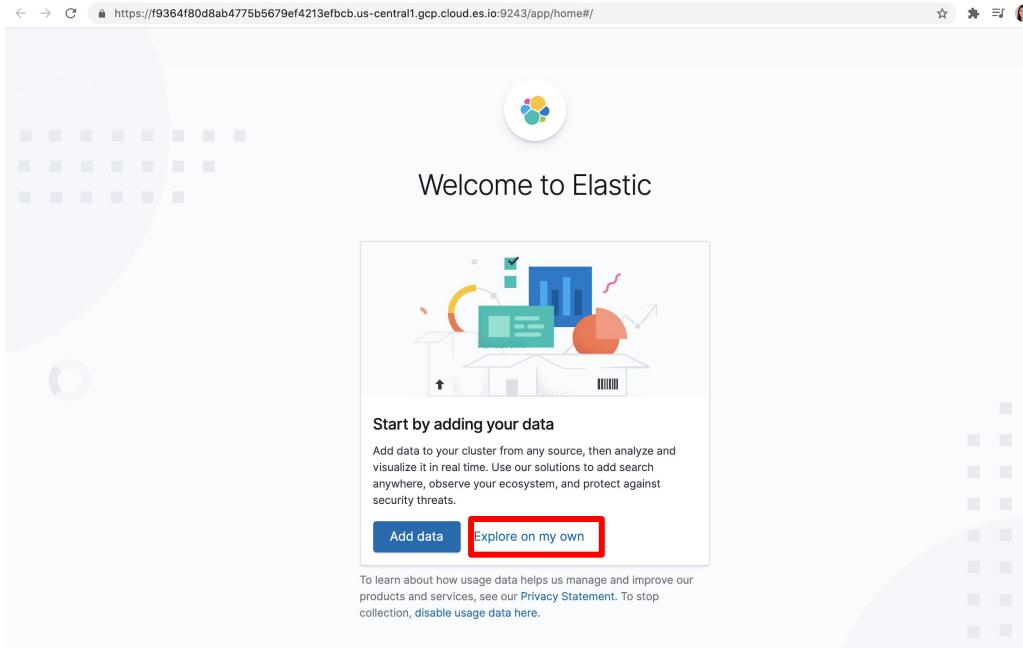
Open Kibana  [Manage](#) ▾

Observability
APM automatically collects logs and metrics from your services and allows you to monitor them.
Logs Import logs from your data sources or logstash and visualize them with preconfigured dashboards.
Metrics Collect metrics from the operating systems and services running on your servers.
Add APM Add log data Add metric data Add events

Visualize and Explore Data
Logs and metrics can be visualized with dashboards, timelines, and search. You can also use the search bar to quickly find specific data.
Metrics Visualize and explore data with charts, maps, and other visualization types.
Logs Use Elasticsearch data to build dashboards and reports.
Add logs Add metrics

Manage and Administer the Elastic Stack
Logs and metrics can be visualized with dashboards, timelines, and search. You can also use the search bar to quickly find specific data.
Metrics Visualize and explore data with charts, maps, and other visualization types.
Logs Use Elasticsearch data to build dashboards and reports.
Add logs Add metrics

Click on Explore on my own option.



Click on Upload a file option

The screenshot shows the Elastic Home page. At the top, there's a navigation bar with the Elastic logo, a search bar, and user icons. Below the navigation is a main content area titled "Home". This area features three cards on the left: "Enterprise Search" (green), "Observability" (pink), and "Security" (dark grey). On the right, there's a large card for "Kibana" (blue) with a sub-section for "Visualize & analyze". At the bottom, there's a section titled "Ingest your data" with three options: "Add data", "Add Elastic Agent", and "Upload a file". The "Upload a file" option is highlighted with a red box. A "Try our sample data" link is located above the "Upload a file" box.

Elastic

Search Elastic

Home

Add data Manage Dev tools

Enterprise Search
Search everything →

Build a powerful search experience.
Connect your users to relevant data.
Unify your team content.

Observability
Centralize & monitor →

Monitor infrastructure metrics.
Trace application requests.
Measure SLAs and react to issues.

Security
SIEM & Endpoint Security →

Prevent threats autonomously.
Detect and respond.
Investigate incidents.

Kibana
Visualize & analyze →

Analyze data in dashboards.
Search and find insights.
Design pixel-perfect presentations.
Plot geographic data.
Model, predict, and detect.
Reveal patterns and relationships.

Ingest your data

Add data
Ingest data from popular apps and services.

Add Elastic Agent
Add and manage your fleet of Elastic Agents and integrations.

Upload a file
Import your own CSV, NDJSON, or log file.

Try our sample data

Download and Unzip News Category Dataset from Kaggle

The screenshot shows the Kaggle website displaying the 'News Category Dataset' by Rishabh Misra. The dataset was updated 2 years ago (Version 2). It has a size of 25 MB and is licensed under CC0: Public Domain. The dataset contains around 200k news headlines from 2012 to 2018 obtained from HuffPost. The Data Explorer shows a file named 'News_Category_Dataset_v2.json' (80.03 MB). The About this file section indicates the file contains 202,372 records and each JSON record includes attributes for category, headline, and authors.

Dataset

News Category Dataset

Identify the type of news based on headlines and short descriptions

Rishabh Misra • updated 2 years ago (Version 2)

Data Tasks Notebooks (64) Discussion (3) Activity Metadata

Download (25 MB) New Notebook

Usability 10.0 License CC0: Public Domain Tags news, nlp, classification, deep learning, linguistics

Description

Context

This dataset contains around 200k news headlines from the year 2012 to 2018 obtained from [HuffPost](#). The model trained on this dataset could be used to identify tags for untracked news articles or to identify the type of language used in different news articles.

Content

Each news headline has a corresponding category. Categories and corresponding article counts are as follows:

- POLITICS: 32739
- ENTERTAINMENT: 17200
- TECHNOLOGY: 15000
- SCIENCE: 14000
- SOCIETY: 13000
- ARTS: 12000
- SPORTS: 11000
- OPINION: 10000
- WORLD: 9000
- BUSINESS: 8000
- CULTURE: 7000
- HEALTH: 6000
- CRIME: 5000
- RELIGION: 4000
- ENVIRONMENT: 3000
- FOOD: 2000
- DRINKS: 1000

Data Explorer

80.03 MB

{ } News_Category_Dataset_v2...

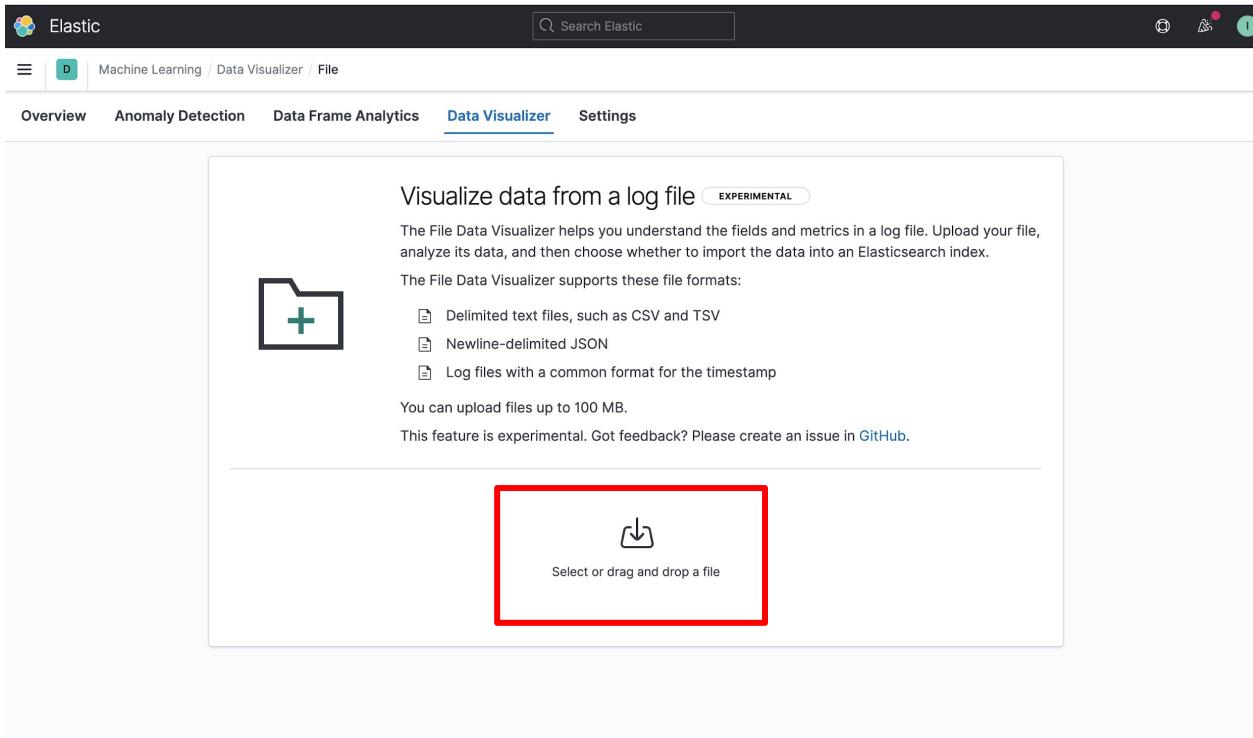
< News_Category_Dataset_v2.json (80.03 MB)

About this file

The file contains 202,372 records. Each json record contains following attributes:

category : Category article belongs to
headline : Headline of the article
authors : Person authored the article

Drag and drop a file you want to upload.



Elastic

Machine Learning / Data Visualizer / File

Overview Anomaly Detection Data Frame Analytics Data Visualizer Settings

Visualize data from a log file EXPERIMENTAL

The File Data Visualizer helps you understand the fields and metrics in a log file. Upload your file, analyze its data, and then choose whether to import the data into an Elasticsearch index.

The File Data Visualizer supports these file formats:

- Delimited text files, such as CSV and TSV
- Newline-delimited JSON
- Log files with a common format for the timestamp

You can upload files up to 100 MB.

This feature is experimental. Got feedback? Please create an issue in [GitHub](#).

Select or drag and drop a file

Kibana will give you an analysis of the first 1000 lines of your data and give you a summary of your dataset.

Elastic

Machine Learning / Data Visualizer / File

Overview Anomaly Detection Data Frame Analytics Data Visualizer Settings

News_Category_Dataset_v2.json

File contents

First 1,000 lines

```
1 {"category": "CRIME", "headline": "There Were 2 Mass Shootings In Texas Last Week, But Only 1 On TV", "authors": "Melissa Jeltsen", "link": "https://www.huffingtonpost.com/entry/texas-amanda-painter-mass-shooting_us_5b081ab4e4b0802d69caad89", "short_description": "She left her husband. He killed their children. Just another day in America." "date": "2018-05-26"}  
2 {"category": "ENTERTAINMENT", "headline": "Will Smith Joins Diplo And Nicky Jam For The 2018 World Cup's Official Song", "authors": "Andy McDonald", "link": "https://www.huffingtonpost.com/entry/will-smith-joins-diplo-and-nicky-jam-for-the-official-2018-world-cup-song_us_5b09726fe4b0fdb2aa541201", "short_description": "Of course it ", "date": "2018-05-26"}  
3 {"category": "ENTERTAINMENT", "headline": "Hugh Grant Marries For The First Time At Age 57", "authors": "Ron Dicker", "link": "https://www.huffingtonpost.com/entry/hugh-marries_us_5b09212c4b0568a880b98c", "short_description": "The actor and his longtime girlfriend Anna Eberstein tied the knot in a civil ceremony.", "date": "2018-05-26"}  
4 {"category": "ENTERTAINMENT", "headline": "Jim Carrey Blasts 'Castrato' Adam Schiff And Democrats In New Artwork", "authors": "Ron Dicker", "link": "https://www.huffingt /entry/jim-carrey-adam-schiff-democrats_us_5b0950e8e4b0fdb2aa53e675", "short_description": "The actor gives Dems an ass-kicking for not fighting hard enough against ", "date": "2018-05-26"}  
5 {"category": "ENTERTAINMENT", "headline": "Julianne Margulies Uses Donald Trump Poop Bags To Pick Up After Her Dog", "authors": "Ron Dicker", "link": "https://www.huffir /entry/julianne-margulies-takes-down-his-trump-poop-bags-to-pick-up-after-her-dog_us_5b0950e8e4b0fdb2aa53e678", "short_description": "The "Nightlands" actress paid union the honor in a "really cathartic" as
```

Summary

Number of lines analyzed 1000

Format ndjson

Time field date

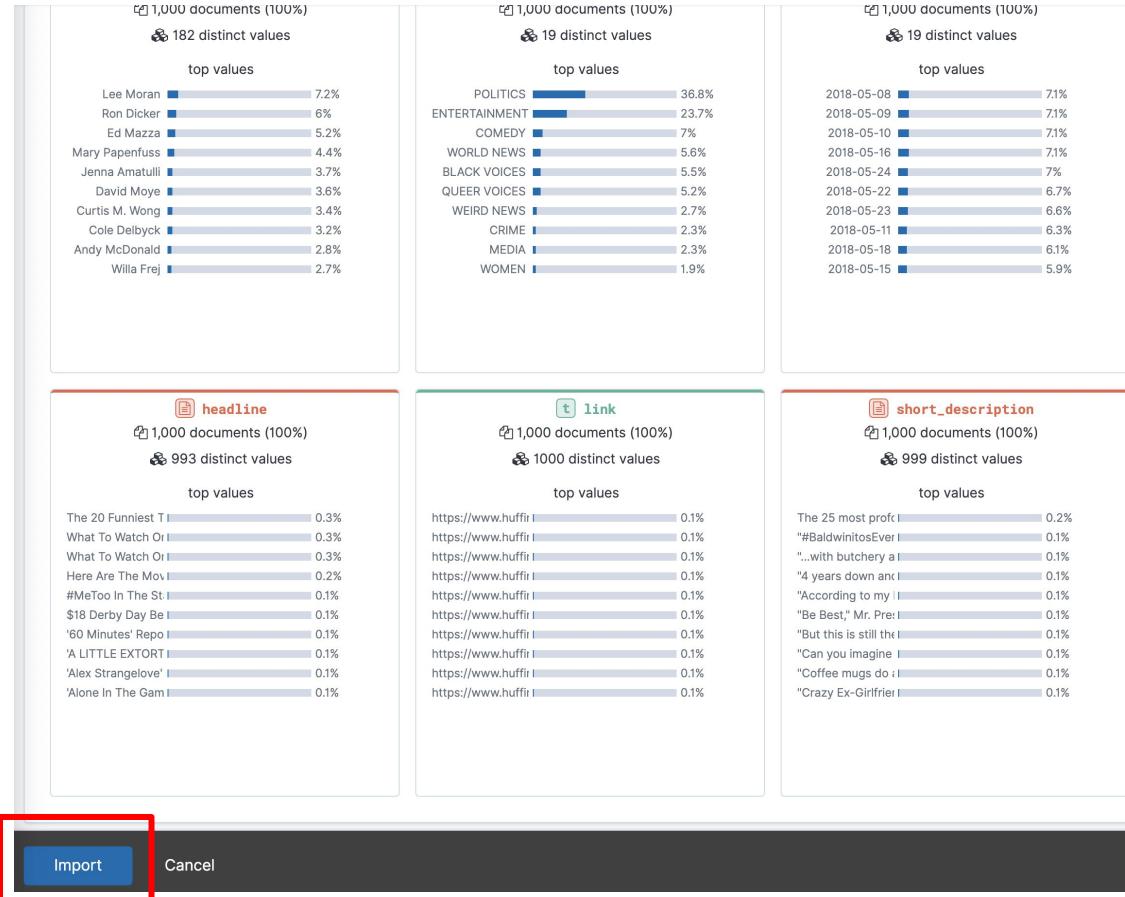
Time format ISO8601

Override settings Analysis explanation

Field section displays fields identified, high level statistics, and top occurring values



Click on import button



Name your index and click on import.

The screenshot shows the Elastic Data Visualizer interface. At the top, there is a dark header bar with the Elastic logo, a search bar containing "Search Elastic", and three small icons on the right. Below the header, the navigation bar includes "Machine Learning / Data Visualizer / File" and tabs for "Overview", "Anomaly Detection", "Data Frame Analytics", "Data Visualizer" (which is underlined in blue), and "Settings".

The main content area displays a file named "News_Category_Dataset_v2.json". Below it is a form titled "Import data" (marked as "EXPERIMENTAL"). The form has two tabs: "Simple" (which is selected) and "Advanced".

The "Index name" field contains "news_headlines". A checkbox labeled "Create index pattern" is checked. At the bottom of the form is a blue "Import" button, which is highlighted with a red rectangular border.

Then Elasticsearch will do the rest!

The screenshot shows the Elasticsearch Data Visualizer interface. At the top, there's a navigation bar with icons for Home, Machine Learning, Data Visualizer, and File. Below the navigation is a search bar labeled "Search Elastic". The main menu includes Overview, Anomaly Detection, Data Frame Analytics, **Data Visualizer** (which is underlined), and Settings. A red box highlights the "Data Visualizer" tab.

The main content area is titled "News_Category_Dataset_v2.json". It features an "Import data" section with an "EXPERIMENTAL" label. Under "Simple", the "Index name" is set to "news_headlines". There's also a checked checkbox for "Create index pattern". A "Reset" button is at the bottom of this section.

Below this is a timeline with five circular markers connected by a horizontal line, each with a checkmark:

- File processed
- Index created
- Ingest pipeline created
- Data uploaded
- Index pattern created

Under the timeline, a section titled "Import complete" shows the following details:

Index	news_articles
Index pattern	news_articles
Ingest pipeline	news_articles-pipeline
Documents ingested	200853

At the bottom of the interface are "Back" and "Cancel" buttons.

Click on menu icon, and open Dev Tools.

The screenshot shows the Elastic Stack interface with the following details:

- Header:** Elastic, Search Elastic, and various icons.
- Left Sidebar:**
 - Home:** Recently viewed (empty), Observability (Overview, Logs, Metrics, APM, Uptime, User Experience), Security (Overview, Detections, Hosts, Network, Timelines, Cases, Administration), and Management (Dev Tools, Fleet, Stack Monitoring, Stack Management). The "Dev Tools" item is highlighted with a blue bar.
- Main Content:** Data Visualizer tab is active, showing a file named "v2.json". Below it is a timeline with four events:
 - Index created
 - Ingest pipeline created
 - Data uploaded
 - Index pattern created
- Bottom:** A footer bar with several icons.

Click on Explore on my own option.

The screenshot shows the Elasticsearch Dev Tools interface. On the left, the 'Console' tab is selected, displaying a search request:

```
GET _search
{
  "query": {
    "match_all": {}
  }
}
```

A blue box highlights this request. On the right, a modal window titled 'Welcome to Console' provides an introduction to the UI:

Quick intro to the UI

The Console UI is split into two panes: an editor pane (left) and a response pane (right). Use the editor to type requests and submit them to Elasticsearch. The results will be displayed in the response pane on the right side.

Console understands requests in a compact format, similar to cURL:

```
1 # index a doc
2 PUT index/_doc/1
3 {
4   "body": "here"
5 }
6
7 # and get it ...
8 GET index/_doc/1
```

While typing a request, Console will make suggestions which you can then accept by hitting Enter/Tab. These suggestions are made based on the request structure as well as your indices and types.

A few quick tips, while I have your attention

- Submit requests to ES using the green triangle button.
- Use the wrench menu for other useful things.
- You can paste requests in cURL format

Dismiss

Fine tuning precision or recall using Elasticsearch and Kibana



Questions?



Join the Elastic Austin User Group for updates on future workshops!

Meetup

Start a new group
50% OFF

Exp



Part of **Elastic – 142 groups**

Elastic Austin User Group

Austin, TX

589 members · Public group

Organized by Elastic Meetup Team and 1 other

Share: [f](#) [t](#) [in](#)

About

Events

Members

Photos

Discussions

More

Manage group

Create event



Lisa Jung

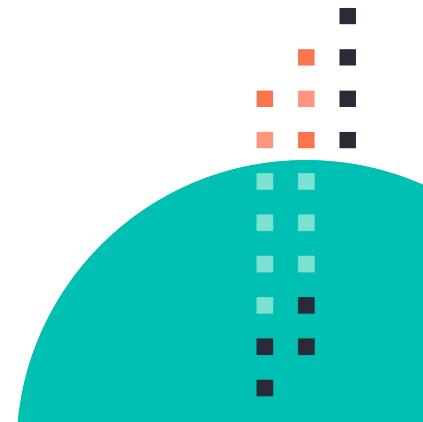
Developer Advocate @Elastic



Discussion forum: <https://discuss.elastic.co/>

Blog: <https://dev.to/lisahjung>

Twitter: @LisaHJung

A large teal circle is positioned in the bottom right corner. To its right, there is a vertical column of small colored squares in red, orange, and black.